

Sparse Autoencoders Can Learn Graded Latents for Relational Composition

Anonymous Authors¹

Abstract

Transformers trained on an ordered bigram copying task can represent multiple features using one set of shared embedding directions, with meaning carried partly by differences in attention-weight magnitudes. This challenges the assumption that transformer features, and sparse autoencoder (SAE) latents trained on them, are independently interpretable. We train SAEs on a transformer known to use this mechanism and find that they reliably learn graded latent activations, where different magnitudes correspond to different list orderings. Steering these latents changes the order of model outputs, suggesting that their magnitudes are causally meaningful. This case study motivates closer attention to activation magnitude when interpreting SAE latents, including in larger pretrained SAEs such as Gemma Scope. Anonymised code is available here: <https://anonymous.4open.science/r/graded-latents-70FA>.

1. Introduction

Transformers can learn mechanisms for relational composition (Wattenberg & Viégas, 2024), in which multiple feature vectors are combined into a single fixed-length representation encoding structured relationships. For example, Farrell et al. (2025) train a transformer on an ordered bigram copying task that learns a relative magnitude-based composition mechanism. The transformer represents both (x, y) and (y, x) as linear combinations of shared embedding directions, with ordering determined by continuously varying coefficients derived from the attention pattern.

This poses an interesting challenge for sparse autoencoders (SAEs) (Bricken et al., 2023; Huben et al., 2024) and their variants (Gao et al., 2025; Bussmann et al., 2024; Rajamanoharan et al., 2024b; Bussmann et al., 2025; Karvo-

nen et al., 2025; Leask et al., 2025b). If the transformer reuses the shared embedding directions while varying their relative magnitudes, then an SAE trained to sparsely reconstruct these activations may recover latents aligned to the underlying embedding directions while encoding ordering information in latent activation magnitude. We refer to such latents as *graded latents*, for which different magnitudes correspond to different compositional states rather than simply indicating whether a feature is present. This challenges interpretation methods that primarily treat SAE latents as feature presence indicators (Paulo et al., 2025; Quirke et al., 2025). This issue persists even if an SAE successfully recovers the underlying feature directions, separating it from previous work focused on dictionary correctness (Chanin et al., 2026; Leask et al., 2025a).

However, Farrell et al. (2025) did not study the impact on SAEs. In this paper, we train a suite of BatchTopK (Bussmann et al., 2024), JumpReLU (Rajamanoharan et al., 2024b) and Matryoshka (Bussmann et al., 2025) SAEs on the toy transformer from Farrell et al. (2025) and find that they reliably learn graded latents whose activation magnitudes correspond to different bigram orderings.

In Section 2 we describe the transformer task, SAE training setup, and graded latent identification and steering protocol. In Section 3, we identify two graded SAE latents using correlations with the transformer’s relative attention variable, $\Delta\alpha$, and show through steering interventions that scaling their activations can causally swap the model output from (x, y) to (y, x) . Across a broad SAE sweep, high-performing SAEs tend to contain latents whose activations correlate with $\Delta\alpha$. For a selected set of six high-performing BatchTopK, JumpReLU, and Matryoshka SAEs, steering these latents swaps between 25–73% of the transformer’s original outputs. We view this work as a preliminary case study motivating further investigation into graded latents in large pretrained SAEs such as Gemma Scope (Lieberum et al., 2024; McDougall et al., 2025).

2. Methodology

We study a two-layer attention-only transformer trained on the ordered bigram copying task introduced by Farrell et al. (2025) (Table 3). The model encodes bigram (d_1, d_2) into a single separator token, s , and reconstructs

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.




Figure 1. This plot shows the logits at output positions o_1 (left) and o_2 (right) for input $(20, 49)$ as latent 11’s activation is scaled by factor p (x-axes). Each line shows the logit for one symbol; blue = d_1 , red = d_2 and all others are grey. A valid swap requires the red line to lead at o_1 and the blue line to lead at o_2 simultaneously. Here, this holds for $p \in [0.16, 2.20]$.

it as $(o_1, o_2) = (d_1, d_2)$. The 5-token input is represented as $[d_1, d_2, s, o_1, o_2]$, where d_1 and d_2 are uniformly sampled from $[0, 99]$ and s, o_1, o_2 are special tokens. This gives 10,000 possible ordered bigrams. Farrell et al. (2025) find that the residual at s after the first layer, $\mathbf{r}_s^{L_1}$, is:

$$\mathbf{r}_s^{L_1} = \alpha_{s \rightarrow d_1}(\mathbf{E}_{d_1} + \mathbf{P}_{d_1}) + \alpha_{s \rightarrow d_2}(\mathbf{E}_{d_2} + \mathbf{P}_{d_2}) + \mathbf{E}_s + \mathbf{P}_s$$

where $\alpha_{s \rightarrow d_i}$ is the attention weight in the first layer for query s (separator token) and key d_i (element of input bigram), and $\mathbf{E}_x$ and $\mathbf{P}_x$ are the token and positional embeddings for the token at x respectively. Thus, the model represents bigram orderings at s using a linear combination of shared embedding directions, $(\mathbf{E}_{d_1} + \mathbf{P}_{d_1})$ and $((\mathbf{E}_{d_2} + \mathbf{P}_{d_2}))$, with coefficients defined by the attention probabilities.

Models can represent more than n features in n -dimensional activations through superposition, which complicates interpretability (Elhage et al., 2022). SAEs decompose dense, superposed model activations into sparse feature representations (Bricken et al., 2023). We train BatchTopK (Bussmann et al., 2024), JumpReLU (Rajamanoharan et al., 2024b) and Matryoshka (Bussmann et al., 2025) SAEs on $\mathbf{r}_s^{L_1}$ using implementations from Marks et al. (2024). SAEs are trained on residual activations on the full finite input space of the toy task (10,000 inputs). Because our goal is not out-of-distribution generalisation, we do not need a train/test split. Our goal is instead to determine whether SAEs recover the known internal computation implemented by the transformer. We evaluate SAEs on sparsity (L_0) and loss recovered (Karvonen et al., 2025; Ferrando et al., 2024) over the complete activation distribution.

Table 1. Selected SAE configuration. The optimiser, betas, schedule and auxiliary loss are fixed in the implementation (Marks et al., 2024) and excluded here for brevity.

Parameter	Value
SAE type	BatchTopK
Input dimension (d_{model})	64
Dictionary size (d_{SAE})	128
k	3
Actual sparsity (L_0)	3.00
Activation centring	Mean-centred (pre-training)
Encoder init	$W_{\text{enc}} = W_{\text{dec}}^\top$
Decoder normalisation	Unit-norm columns
Learning rate	3×10^{-4}
Batch size	4096
Training steps	50,000
Seed	1
Hardware	NVIDIA GeForce RTX 2080 Ti

After a grid search (Table 5) we find that BatchTopK SAEs with $k \in [3, 5]$ all achieve similar top performance, so we select $k = 3$. We select $d_{\text{SAE}} = 128$ which achieves similar performance to larger d_{SAE} . We filter for loss recovered over 0.999 and use the checkpoint with the fewest dead latents (3.1%). Table 1 shows the final configuration.

We inspect latent activations using top-activating examples, firing rates, and correlations with the difference between transformer attention weights, $\Delta\alpha = \alpha_{s \rightarrow d_1} - \alpha_{s \rightarrow d_2}$, because Farrell et al. (2025) find $\Delta\alpha$ differentiates bigram orderings at $\mathbf{r}_s^{L_1}$. We hypothesise that latents related to or-

dering will correlate with $\Delta\alpha$ significantly more strongly than latents indicating feature presence. To test whether a latent is graded, we perform causal steering interventions inspired by O’Brien et al. (2025). For some latent, we replace its activation magnitude m with $m' = pm$ for some scale factor p . We reconstruct and replace the original activation at $r_s^{L_1}$ and observe changes in the output logits. We repeat across all $p \in [0, 10]$ (step 0.05) and all input bigrams. If the latent is graded, we expect different scale factors to systematically change the model’s output ordering.

We test for latents correlated with $\Delta\alpha$ across a sweep of 2,613 SAEs spanning multiple architectures and hyperparameters (Table 4). For computational feasibility, we only perform steering on six selected high-performing checkpoints (Table 8) chosen to span architecture, dictionary size and sparsity settings. We therefore treat the steering results as exploratory rather than statistically significant.

3. Results

3.1. Analysing Features

We visualise latent activations over all 100×100 ordered bigrams (d_1, d_2) using heatmaps (Figure 2). We identify two categories, excluding 8% of latents which activate on fewer than 0.1% of inputs: *k-symbol detectors* (91% of latents) and *special latents* (2 latents (1.6%)). 88% are 1-symbol detectors: they activate only if one specific symbol is present at either d_1 or d_2 , producing a ‘plus-sign’ pattern in Figure 2 and activating most strongly when $d_1 = d_2$. 3% of latents detect multiple symbols ($k > 1$) simultaneously, showing multiple ‘plus-signs’. Only two latents (11 and 56) do not follow the above pattern. Latent 11 activates on 64% of inputs and latent 56 on 9.9%, compared to fewer than 4.5% for all other latents (Appendix C). Because they activate across most inputs, they cannot be presence indicators for specific symbols.

Moreover, Latent 11 strongly positively correlates with $\Delta\alpha$ (Pearson $r = 0.807$) over the whole dataset, while latent 56 has a weaker negative correlation ($r = -0.3213$). All other latents have $|r| < 0.004$. This sharply distinguishes latents 11 and 56 from others and suggests that they may encode ordering information. We use $|r| = 0.3$ as a threshold for identifying special latents because of this significant empirical gap with others. Since $\Delta\alpha$ is positive when attention is weighted towards d_1 , we refer to latent 11 as strongly d_1 -favouring and latent 56 as weakly d_2 -favouring (Appendix C, Figure 4).

3.2. Steering Experiments

Steering latent 11 swaps for 50% of model outputs (Figure 1). There are several failure reasons (visualisations in Appendix D): **No activation (36%)**: Some inputs do not

activate latent 11 ($m = 0$). **Negative scaling required (9.2%)**: Swapping o_1 requires negative p , which we reject as BatchTopK activations are non-negative by construction. **Mutually exclusive swap ranges (2.3%)**: Each output position can be individually swapped but not simultaneously. **No swap in observed range (1.8%)**: One logit remains dominant across all tested p . **Matching bigram elements (0.8%)**: $d_1 = d_2$, making the swap degenerate. However, it is sometimes possible to scale a co-activating (non-special) latent to change the logit distribution and successfully swap more outputs with latent 11, so 50% is not the ceiling (see Appendix D). Steering latent 56 swaps 2.0% of outputs and never coactivates with latent 11, suggesting they may be complementary. We additionally test steering non-special latents and do not find comparable behaviour (Appendix D), which suggests that it is specific to special latents.

Across inputs, logits at o_1 scale approximately linearly with p (Figure 1) and linear regression achieves mean $R^2 > 0.999$. We therefore estimate exact swap thresholds at o_1 analytically using lines fitted to the d_1 and d_2 logits.

3.3. Generalisation to other SAEs

Across 2,613 SAEs, high-performing ones frequently contain special latents, which we identify using $\Delta\alpha$ correlation $|r| > 0.3$ (Appendix E, Figure 12).

We implement a pipeline based on empirical findings to identify special latents and steer to swap model outputs in six selected high-performing checkpoints (see Appendix B, Table 8) All six learned at least one special latent and five learned a pair of opposing special latents, one d_1 -favouring and one d_2 -favouring. Steering at least one special latent swapped a substantial proportion of outputs, ranging from 26.6% to 72.5% of the dataset (Table 10). Conditional on latent activation, these were substantially higher. Swap failures primarily occur for the same reasons as in Section 3.2. This suggests that steerable graded latents are not unique to a single BatchTopK checkpoint, although the limited subset of evaluated SAEs limits statistical significance. Further results are in Appendix E.

4. Discussion

Our results suggest a failure mode for interpreting SAE latents as binary feature presence indicators. The k -symbol detectors in Figure 2 behave as binary indicators. In contrast, the special latents are difficult to interpret from support alone. Latent 11 activates on a broad subset of inputs, and its activation magnitude tracks $\Delta\alpha$, a continuous variable in the transformer’s attention pattern. Different magnitudes then correspond to different downstream ordering behavior. Thus, the latent is not merely ‘on’ for an ordering feature; its scalar value participates in the ordering computation.




Figure 2. These 100×100 heatmaps show example SAE latent activations over all input bigrams (d_1, d_2) with $d_1, d_2 \in [0, 99]$: top-left cell is $(0, 0)$, bottom-right cell is $(99, 99)$. **Left** (1-symbol detector): latent 0 activates strongly if d_1 or $d_2 = 41$ and peaks at $d_1 = d_2 = 41$, producing a plus-sign pattern. **Centre** (k -symbol detector, $k > 1$): latent 5 activates strongly if d_1 or $d_2 \in \{7, 58\}$, producing two plus-sign patterns. **Right** (special latent): latent 11 activates on most inputs with generally much greater magnitudes than any other latent.

This suggests that SAE interpretability tools should not only ask which latents activate, but also whether activation magnitudes encode causal computational variables.

However, this toy transformer is limited ...

TODO:

Explain the asymmetry between paired special latents

In the primary SAE, latent 11 has $r=0.807$ and swaps about 50

This is interesting and currently under-discussed. Possible explanations include top-k competition, asymmetric transformer attention patterns, uneven activation support, or SAE training artifacts. Even a short paragraph saying “we do not yet understand this asymmetry” would be useful.

5. Conclusion and Future Work

In this paper, we showed that SAEs trained on an activation that uses relative magnitude-based composition reliably learn latents with graded activations with magnitude that correlate with $\Delta\alpha$. We demonstrated that steering these latents predictably changes the transformer’s predicted output. We release our trained SAEs on Weights & Biases to support future work and reproducibility of our results. Limitations are listed in Appendix F.

In future work we will attempt to find graded latent activations in real-world pretrained SAEs. The Gemma Scope releases (Lieberum et al., 2024; McDougall et al., 2025) provide a natural test environment. They include JumpReLU and Matryoshka SAEs trained on Gemma 2 and Gemma 3 at multiple widths, matching the SAE variants studied here. The specific target would be to identify SAE latents in pretrained models where activation magnitude, rather than feature presence, is causally relevant for downstream

behaviour.

In summary, the findings in this paper are intended as a foundation for further research on magnitude-based feature composition mechanisms in LLMs.

Impact Statement

This paper presents work whose goal is to advance our understanding of sparse autoencoders, which are popular tools in interpretability research. If our findings are also present in real-world language models, then this may present significant problems for use of SAEs for interpretability of such systems.

References

- Bricken, T., Templeton, A., Batson, J., Chen, B., Jermyn, A., Conerly, T., Turner, N., Anil, C., Denison, C., Askell, A., Lasenby, R., Wu, Y., Kravec, S., Schiefer, N., Maxwell, T., Joseph, N., Hatfield-Dodds, Z., Tamkin, A., Nguyen, K., McLean, B., Burke, J. E., Hume, T., Carter, S., Henighan, T., and Olah, C. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023. <https://transformer-circuits.pub/2023/monosemantic-features/index.html>.
- Brix, G., Durrant, M. G., Ku, J., Naghipourfar, M., Poli, M., Sun, G., Brockman, G., Chang, D., Fanton, A., Gonzalez, G. A., King, S. H., Li, D. B., Merchant, A. T., Nguyen, E., Ricci-Tam, C., Romero, D. W., Schmok, J. C., Taghibakhshi, A., Vorontsov, A., Yang, B., Deng, M., Gorton, L., Nguyen, N., Wang, N. K., Pearce, M. T., Simon, E., Adams, E., Amador, Z. J., Ashley, E. A., Bacchus, S. A., Dai, H., Dillmann, S., Ermon, S., Guo, D., Herschl, M. H., Ilango, R., Janik, K., Lu, A. X., Mehta, R., Mofrad, M. R. K., Ng, M. Y., Pannu, J., Ré, C., St. John,

- J., Sullivan, J., Tey, J., Viggiano, B., Zhu, K., Zynda, G., Balsam, D., Collison, P., Costa, A. B., Hernandez-Boussard, T., Ho, E., Liu, M.-Y., McGrath, T., Powell, K., Pinglay, S., Burke, D. P., Goodarzi, H., Hsu, P. D., and Hie, B. L. Genome modelling and design across all domains of life with evo 2. *Nature*, 03 2026. ISSN 1476-4687. doi: 10.1038/s41586-026-10176-5. URL <https://doi.org/10.1038/s41586-026-10176-5>.
- Busmann, B., Leask, P., and Nanda, N. Batchtopk sparse autoencoders. In *NeurIPS 2024 Workshop on Scientific Methods for Understanding Deep Learning*, 2024. URL <https://openreview.net/forum?id=d4dpOCqybL>.
- Busmann, B., Nabeshima, N., Karvonen, A., and Nanda, N. Learning multi-level features with matryoshka sparse autoencoders. In *Forty-second International Conference on Machine Learning*, 2025. URL <https://openreview.net/forum?id=m25T5rAy43>.
- Chanin, D., Wilken-Smith, J., Dulka, T., Bhatnagar, H., Golechha, S., and Bloom, J. I. A is for absorption: Studying feature splitting and absorption in sparse autoencoders. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2026. URL <https://openreview.net/forum?id=R73ybUciQF>.
- Elhage, N., Hume, T., Olsson, C., Schiefer, N., Henighan, T., Kravec, S., Hatfield-Dodds, Z., Lasenby, R., Drain, D., Chen, C., et al. Toy models of superposition. *arXiv preprint arXiv:2209.10652*, 2022.
- Farrell, T., Leask, P., and Moubayed, N. A. Order by scale: Relative-magnitude relational composition in attention-only transformers. In *Socially Responsible and Trustworthy Foundation Models at NeurIPS 2025*, 2025. URL <https://openreview.net/forum?id=vWRVzNtk7W>.
- Ferrando, J., Sarti, G., Bisazza, A., and Costa-Jussà, M. R. A primer on the inner workings of transformer-based language models. *arXiv preprint arXiv:2405.00208*, 2024.
- Gao, L., la Tour, T. D., Tillman, H., Goh, G., Troll, R., Radford, A., Sutskever, I., Leike, J., and Wu, J. Scaling and evaluating sparse autoencoders. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=tcsZt9ZNKD>.
- Huben, R., Cunningham, H., Smith, L. R., Ewart, A., and Sharkey, L. Sparse autoencoders find highly interpretable features in language models. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=F76bwRSLeK>.
- Karvonen, A., Rager, C., Lin, J., Tigges, C., Bloom, J. I., Chanin, D., Lau, Y.-T., Farrell, E., McDougall, C. S., Ayonrinde, K., Till, D., Wearden, M., Conmy, A., Marks, S., and Nanda, N. SAE Bench: A comprehensive benchmark for sparse autoencoders in language model interpretability. In *Forty-second International Conference on Machine Learning*, 2025. URL <https://openreview.net/forum?id=qrU3yNfX0d>.
- Leask, P., Busmann, B., Pearce, M. T., Bloom, J. I., Tigges, C., Moubayed, N. A., Sharkey, L., and Nanda, N. Sparse autoencoders do not find canonical units of analysis. In *The Thirteenth International Conference on Learning Representations*, 2025a. URL <https://openreview.net/forum?id=9ca9eHNrdH>.
- Leask, P., Nanda, N., and Moubayed, N. A. Inference-time decomposition of activations (ITDA): A scalable approach to interpreting large language models. In *Forty-second International Conference on Machine Learning*, 2025b. URL <https://openreview.net/forum?id=4WMZqAoYi4>.
- Lieberum, T., Rajamanoharan, S., Conmy, A., Smith, L., Sonnerat, N., Varma, V., Kramar, J., Dragan, A., Shah, R., and Nanda, N. Gemma scope: Open sparse autoencoders everywhere all at once on gemma 2. In Belinkov, Y., Kim, N., Jumelet, J., Mohebbi, H., Mueller, A., and Chen, H. (eds.), *Proceedings of the 7th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, pp. 278–300, Miami, Florida, US, November 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.blackboxnlp-1.19. URL <https://aclanthology.org/2024.blackboxnlp-1.19/>.
- Marks, S., Karvonen, A., and Mueller, A. dictionary.learning, 2024. URL https://github.com/sapmarks/dictionary_learning.
- McDougall, C., Conmy, A., Kramár, J., Lieberum, T., Rajamanoharan, S., and Nanda, N. Gemma scope 2: Technical paper. Technical report, Google DeepMind, 9 2025. URL https://storage.googleapis.com/deepmind-media/DeepMind.com/Blog/gemma-scope-2-helping-the-ai-safety-community-deepen-understanding-of-complex-language-model-behavior/Gemma_Scope_2_Technical_Paper.pdf.
- O’Brien, K., Majercak, D., Fernandes, X., Edgar, R. G., Bullwinkel, B., Chen, J., Nori, H., Carignan, D., Horvitz, E., and Poursabzi-Sangdeh, F. Steering language model refusal with sparse autoencoders. In *ICML 2025 Workshop on Reliable and Responsible Foundation Models*, 2025. URL <https://openreview.net/forum?id=PMK1jdGQoc>.

- Paulo, G. S., Mallen, A. T., Juang, C., and Belrose, N. Automatically interpreting millions of features in large language models. In *Forty-second International Conference on Machine Learning*, 2025. URL <https://openreview.net/forum?id=EemtbbhJOXc>.
- Quirke, L., Shabalin, S., and Belrose, N. Binary sparse coding for interpretability. *arXiv preprint arXiv:2509.25596*, 2025.
- Rajamanoharan, S., Conmy, A., Smith, L., Lieberum, T., Varma, V., Kramar, J., Shah, R., and Nanda, N. Improving sparse decomposition of language model activations with gated sparse autoencoders. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024a. URL <https://openreview.net/forum?id=zLBlin2zvW>.
- Rajamanoharan, S., Lieberum, T., Sonnerat, N., Conmy, A., Varma, V., Kramár, J., and Nanda, N. Jumping ahead: Improving reconstruction fidelity with jumprelu sparse autoencoders. *arXiv preprint arXiv:2407.14435*, 2024b.
- Wattenberg, M. and Viégas, F. Relational composition in neural networks: A survey and call to action. In *ICML 2024 Workshop on Mechanistic Interpretability*, 2024. URL <https://openreview.net/forum?id=zzCEiUIPk9>.

A. Further comparison of SAE Types

Table 2. Comparison of some sparse autoencoder (SAE) variants.

Type	Description
Standard ReLU (Bricken et al., 2023)	Applies an L_1 penalty to hidden activations using a standard ReLU activation function to encourage sparsity.
Gated (Rajamanoharan et al., 2024a)	Decouples feature detection from magnitude estimation using a gated mechanism to prevent the shrinkage issues caused by standard L_1 penalties.
JumpReLU (Rajamanoharan et al., 2024b)	Uses a discontinuous step function at a learned per-feature threshold, approximating an L_0 penalty through straight-through estimation. Notably used in Lieberum et al. (2024).
BatchTopK (Bussmann et al., 2024)	Retains top $n \times k$ activations over a batch of n inputs and sets the rest to zero, where k is a tuned hyperparameter. This forces an average of k latent activations per input, so the actual L_0 sparsity is approximately k . Notably used in Brixi et al. (2026); Karvonen et al. (2025).
Matryoshka (Bussmann et al., 2025)	Trains nested sub-dictionaries of increasing size that share a single encoder/decoder, producing multiple resolutions of feature granularity in one run. Notably used in McDougall et al. (2025).

An SAE consists of a single hidden layer trained to reconstruct model activations subject to a sparsity penalty on the latent activations (typically L_1 or top- k regularisation) (Huben et al., 2024; Bricken et al., 2023). Many SAE variants exist in the literature, some of which are shown in Table 2.

We use 3 types of SAEs in this paper. JumpReLU SAEs (Rajamanoharan et al., 2024b) use a discontinuous step function at a learned threshold and optimise an L_0 penalty. BatchTopK SAEs (Bussmann et al., 2024) keep the top $n \times k$ activations over a batch of n inputs and sets the rest to zero. This forces an average of k latent activations per input, so the actual L_0 sparsity is approximately k . Matryoshka SAEs (Bussmann et al., 2025) train nested sub-dictionaries of increasing size under a shared encoder and decoder, producing representations at multiple levels of granularity from a single training run.

B. Experimental Setup

Table 3. Model and training configuration from Farrell et al. (2025). The model achieved 91.45% test accuracy using an 80/20 train/test split.

Parameter	Value
Number of layers	2
Number of heads	1
Residual stream dimension	64
Attention head dimension	64
LayerNorm	None
Bias	None
Value matrix (W_V)	Identity (frozen)
Output matrix (W_O)	Identity (frozen)
Learning rate	1×10^{-3}
Optimiser	AdamW
Batch size	128
Betas	(0.9, 0.999)
Weight decay	0.01

Table 4. SAE hyperparameter sweep configurations. Braces denote swept values; fixed values are listed directly. Fixed parameters across all runs: 50,000 training steps, 1,000 warmup steps, 4,096 batch size. The sweeps are ongoing, so the counts below report evaluated checkpoints rather than completed grid sizes. Note that $n_{\text{groups}} = 1$ in a Matryoshka SAE reduces it to a standard BatchTopK SAE, so this serves as an internal consistency check.

Parameter	Architecture	Values
d_{SAE}	All	{128, 192, 256, 320}
k / Target L_0	BatchTopK, Matryoshka	{1, 2, 3, 4, 5}
Target L_0	JumpReLU	{1, 2, 3, 4, 5}
Learning rate	BatchTopK, Matryoshka	3×10^{-4}
	JumpReLU	$\{7 \times 10^{-5}, 3 \times 10^{-4}\}$
Sparsity penalty	JumpReLU	{0.1, 0.5, 1.0, 5.0}
n_{groups}	Matryoshka	{1, 2, 3, 4}
Architecture	Seeds	Total runs
BatchTopK	18	492
JumpReLU	15	1,201
Matryoshka	15	920

Table 5. BatchTopK SAE sweep results aggregated over multiple seeds (mean $\pm$ std). L_0 : mean active features per token. d_{SAE} : dictionary size. **Loss Recovered**: fraction of model loss recovered by the SAE reconstruction ($\uparrow$ better). **Dead %**: percentage of dictionary features with zero activation across the evaluation set ($\downarrow$ better). **Bold**: best within each L_0 block; **underlined**: best overall.

L_0	d_{SAE}	Runs	Loss Recovered ($\uparrow$)	Dead % ($\downarrow$)
1	128	30	0.9392 ± 0.0020	12.4 ± 2.5
1	192	30	0.9410 ± 0.0015	43.6 ± 1.6
1	256	24	0.9405 ± 0.0018	54.3 ± 2.3
1	320	15	0.9397 ± 0.0015	61.9 ± 1.3
2	128	28	0.9816 ± 0.0021	17.5 ± 1.6
2	192	29	0.9812 ± 0.0025	42.3 ± 1.3
2	256	23	0.9815 ± 0.0020	54.6 ± 1.5
2	320	15	0.9802 ± 0.0025	61.6 ± 1.8
3	128	31	0.9966 ± 0.0122	14.2 ± 2.8
3	192	30	0.9979 ± 0.0091	35.7 ± 5.0
3	256	19	0.9963 ± 0.0145	47.7 ± 6.1
3	320	13	0.9996 ± 0.0001	54.1 ± 6.9
4	128	31	0.9996 ± 0.0000	1.6 ± 1.9
4	192	23	0.9996 ± 0.0000	5.0 ± 1.9
4	256	24	0.9996 ± 0.0001	19.0 ± 15.9
4	320	16	0.9996 ± 0.0001	19.4 ± 12.6
5	128	30	0.9996 ± 0.0000	0.3 ± 0.5
5	192	38	0.9996 ± 0.0000	4.4 ± 6.0
5	256	27	0.9996 ± 0.0000	7.6 ± 3.2
5	320	15	0.9997 ± 0.0000	11.8 ± 3.2

Table 6. JumpReLU SAE sweep results aggregated over multiple seeds (mean $\pm$ std). L_0 : mean active features per token. d_{SAE} : dictionary size. **Loss Recovered**: fraction of model loss recovered by the SAE reconstruction ($\uparrow$ better). **Dead %**: percentage of dictionary features with zero activation across the evaluation set ($\downarrow$ better). **Bold**: best within each L_0 block; **underlined**: best overall.

L_0	d_{SAE}	Runs	Loss Recovered ($\uparrow$)	Dead % ($\downarrow$)
2	128	15	0.9940 ± 0.0017	19.3 ± 0.7
2	192	5	0.9918 ± 0.0000	46.9 ± 0.0
2	256	13	0.9940 ± 0.0025	56.7 ± 2.6
2	320	5	0.9907 ± 0.0000	67.8 ± 0.0
3	128	333	0.9984 ± 0.0020	18.4 ± 1.2
3	192	320	0.9984 ± 0.0020	42.1 ± 1.8
3	256	305	0.9985 ± 0.0020	52.0 ± 3.4
3	320	148	0.9973 ± 0.0025	60.1 ± 4.8
4	192	35	0.9984 ± 0.0022	40.5 ± 3.3
4	256	15	0.9967 ± 0.0023	55.2 ± 3.8
4	320	7	0.9958 ± 0.0018	62.6 ± 3.9

Table 7. Matryoshka SAE sweep results aggregated over multiple seeds (mean $\pm$ std). L_0 : mean active features per token. d_{SAE} : dictionary size. **Loss Recovered**: fraction of model loss recovered by the SAE reconstruction ($\uparrow$ better). **Dead %**: percentage of dictionary features with zero activation across the evaluation set ($\downarrow$ better). **Bold**: best within each L_0 block; **underlined**: best overall.

L_0	d_{SAE}	Runs	Loss Recovered ($\uparrow$)	Dead % ($\downarrow$)
1	128	45	0.9353 ± 0.0019	13.7 ± 1.7
1	192	45	0.9362 ± 0.0017	42.5 ± 1.3
1	256	36	0.9373 ± 0.0015	53.8 ± 1.6
2	128	45	0.9710 ± 0.0036	15.3 ± 2.9
2	192	45	0.9721 ± 0.0034	41.8 ± 2.6
2	256	35	0.9711 ± 0.0033	54.5 ± 3.4
3	128	58	0.9989 ± 0.0005	12.0 ± 3.0
3	192	59	0.9993 ± 0.0003	37.6 ± 5.3
3	256	57	0.9995 ± 0.0002	46.9 ± 7.2
3	320	40	0.9995 ± 0.0001	50.0 ± 9.5
4	128	61	0.9992 ± 0.0003	2.8 ± 2.6
4	192	58	0.9995 ± 0.0001	15.0 ± 9.2
4	256	61	0.9996 ± 0.0001	13.6 ± 11.7
4	320	48	0.9996 ± 0.0001	15.1 ± 13.3
5	128	61	0.9992 ± 0.0003	0.6 ± 0.6
5	192	61	0.9995 ± 0.0001	8.0 ± 6.0
5	256	61	0.9996 ± 0.0001	5.9 ± 3.8
5	320	44	0.9996 ± 0.0001	6.6 ± 4.9

Table 8. Comparison of selected high-performing SAE checkpoints across configuration and performance metrics. These checkpoints were chosen for the steering analysis to span architecture, dictionary size and sparsity settings while keeping reconstruction performance high. Dashes indicate parameters not applicable to a given architecture. L_0 (actual) is the mean number of active latents measured at eval time. **Bold**: best value per performance metric.

SAE Type	Configuration					Performance			
	d_{SAE}	LR	k / Target L_0	Sp. Pen.	n_{groups}	CE Increase $\downarrow$	Expl. Var. $\uparrow$	Dead (%) $\downarrow$	L_0 (actual)
BatchTopK	192	3×10^{-4}	3	—	—	0.0257	0.9929	26.04	3.36
	128	3×10^{-4}	5	—	—	0.0292	0.9939	0.78	4.90
JumpReLU	256	7×10^{-5}	5	5	—	0.0141	0.9965	50.78	4.03
	128	7×10^{-5}	5	5	—	0.0176	0.9963	17.97	3.24
Matryoshka	256	3×10^{-4}	3	—	2	0.0262	0.9928	39.06	3.32
	192	3×10^{-4}	4	—	4	0.0325	0.9927	12.50	3.98

C. Further SAE Latent Analysis




Figure 3. This histogram shows the distribution of firing rates for the SAE checkpoint’s latents. The y-axis uses a symmetric logarithmic scale with a linear scale between $[-1, 1]$. A given latent’s firing rate is the proportion of inputs which cause that latent to activate. The special latents are the only ones with a firing rate greater than one standard deviation from the mean (mean: 0.0236, standard deviation: 0.0559): latent 11 has rate 0.6400, latent 56 has rate 0.0991.




Figure 4. Left: This plot shows the (Pearson) correlation of each latent with $\Delta\alpha = \alpha_{s \rightarrow d_1} - \alpha_{s \rightarrow d_2}$, which is intuitively the emphasis given to d_1 over d_2 by SEP in the transformer model’s first layer’s attention pattern. The y-axis uses a symmetric logarithmic scale with a linear scale between $[-0.05, 0.05]$. d_1 -favouring latents are highlighted red and d_2 -favouring latents are highlighted blue. The top three strongest correlations are: latent 11 with $r = 0.807$, latent 56 with $r = -0.3213$, and latent 96 with $r = 0.0036$. **Right:** This plot shows the correlation of latent 11 with $\Delta\alpha$ in more detail, as this latent has the strongest correlation. All 10,000 possible input bigrams are plotted, and highlighted blue if they activate latent 11 and grey otherwise. We see that the latent activates with greater magnitude for larger $\Delta\alpha$.

D. Further Latent Steering Examples

Figure 5 shows an example where swapping o_1 requires negative scaling ($p < 0$). The logit for d_2 at o_1 only overtakes d_1 when the latent activation is suppressed below zero, which is outside the valid activation range of a BatchTopK SAE.

Figure 6 shows an example where the swap ranges for o_1 and o_2 are mutually exclusive: there exists a positive p that swaps each position individually, but no single p swaps both simultaneously.

Figures 7 and 8 show examples where no swap occurs across the full observed range $p \in [0, 10]$. In Figure 7, the logit for d_1 remains dominant at o_1 for all p ; Figure 8 shows the analogous failure at o_2 .

In limited experiments with inputs that fail for the first three causes, it is sometimes possible to scale a co-activating (non-special) latent to successfully swap outputs. Figures 9, 10 and 11 show an example where input (41, 49) activates latents 11 (magnitude 3.1), 0 (magnitude 3.1) and 55 (magnitude 2.5). The swap fails with latent 11 because there are no positive overlapping scale ranges; however, scaling latent 55 by $p \in [1.1, 1.15]$ successfully produces a swap. This does not necessarily indicate that latent 55 is graded; rather, scaling it alters the logit distribution, which opens a valid swap zone for latent 11. Similarly, running the steering pipeline on latent 55 alone swaps only 4 inputs, all of which co-activate latent 11, suggesting a similar mechanism. We note this preliminary experiment is insufficient for statistical significance, but it suggests that scaling multiple latents together may enable more output swaps than steering special latents alone.

We do not attempt steering on all latents due to computational constraints, but we attempt it on five randomly selected latents (0, 55, 63, 82, 117) and latent 96 (non-special latent with highest $\Delta\alpha$ correlation, $r = 0.0036$) and latent 5 (which is shown in Figure 2). Results are shown in Table 9. Every successful swap using a non-special latent has a co-activating special latent. This suggests that these non-special latents are unlikely to be graded. Instead, they may alter the logit distribution so that graded special latents can successfully swap outputs.

D.1. Steering Across Multiple SAEs




Figure 5. For 9.2% of inputs, latent 11 (in the BatchTopK SAE from Section 2) must be negatively scaled to swap the predicted symbol at one of the output positions.




Figure 6. For 2.3% of inputs, it is possible to scale latent 11 (in the BatchTopK SAE from Section 2) to swap the symbol at each output position individually but not simultaneously.




Figure 7. For 1.6% of inputs, the logit for the symbol at d_2 is always dominant at o_2 in the observed latent scale range (in the BatchTopK SAE from Section 2), so there is no scale where the the symbol at d_1 is predicted instead.




Figure 8. For 0.2% of inputs, the logit for the symbol at d_1 is always dominant at o_1 in the observed latent scale range (in the BatchTopK SAE from Section 2), so there is no scale where the the symbol at d_2 is predicted instead.




Figure 9. In the BatchTopK SAE from Section 2, latent 11 cannot be scaled to swap input bigram (41, 49) because the logit for 49 is never argmax in position o_1 , as shown by the top plot.




Figure 10. However, in the same SAE, latent 55 co-activates with latent 11 for input (41, 49) and can be scaled by factor $p \in [1.07, 1.18]$ to cause the model to output (49, 41) instead.




Figure 11. Moreover, in the same SAE, if latent 55's activation magnitude is scaled by a factor of 1.1, then we can now scale latent 11 by factor $p \in [0.03, 1.45]$ to swap the outputs, which was not possible without scaling latent 55. This suggests that some graded activation demonstrations are only possible when multiple active latents are scaled simultaneously.

Table 9. We perform steering experiments on several chosen and randomly sampled SAE latents to demonstrate how steering with special latents (**bold**) compares to non-special latents. We record how many outputs are successfully swapped by steering only that latent, how many do not activate the latent, and how many activate the latent but cannot be swapped. We additionally report the proportion of the activating inputs that are swapped (Cond. Swap).

Latent	Swapped	No Activation	Active & No Swap	Cond. Swap
0	6 (0.1%)	9786 (97.9%)	208 (2.1%)	2.8%
5	23 (0.2%)	9594 (95.9%)	383 (3.8%)	5.7%
11	4976 (49.8%)	3600 (36.0%)	1424 (14.2%)	77.8%
55	4 (0.0%)	9787 (97.9%)	209 (2.1%)	1.9%
56	196 (2.0%)	9009 (90.1%)	795 (7.9%)	19.8%
63	4 (0.0%)	9801 (98.0%)	195 (2.0%)	2.0%
82	4 (0.0%)	9787 (97.9%)	209 (2.1%)	1.9%
117	37 (0.4%)	9782 (97.8%)	181 (1.8%)	17.0%

Table 10. Generalisation of special latent steering to the selected SAE checkpoints from Table 8. For each identified special latent ($|r| > 0.3$), we report its bias, Pearson correlation r with $\Delta\alpha$, and the percentage of inputs where the latent was successfully steered to swap outputs (Success), did not activate (No Act.), or activated but could not be swapped (Active & No Swap). Cond. Success is success conditional on the latent activating: $\text{Success}/(\text{Success} + \text{Active No Swap})$.

SAE Type	d_{SAE}	L_0	Latent	Type	r	Steering Outcomes (%)			
						Success	No Act.	Active No Swap	Cond. Success
BatchTopK (Studied)	128	3.00	11	d_1 -favouring	+0.807	49.8%	36.0%	14.2%	77.8%
			56	d_2 -favouring	-0.321	2.0%	90.1%	7.9%	20.2%
BatchTopK	192	3.36	47	d_2 -favouring	-0.836	52.9%	34.4%	12.7%	80.6%
			116	d_1 -favouring	+0.365	0.5%	91.7%	7.8%	6.0%
	128	4.90	68	d_1 -favouring	+0.845	32.9%	44.7%	22.4%	59.5%
			64	d_2 -favouring	-0.730	6.0%	62.0%	32.0%	15.8%
JumpReLU	256	4.03	195	d_1 -favouring	+0.849	38.8%	39.8%	21.4%	64.5%
			179	d_2 -favouring	-0.721	5.8%	60.2%	34.0%	14.6%
	128	3.24	105	d_2 -favouring	-0.816	28.6%	37.8%	33.6%	46.0%
			4	d_1 -favouring	+0.756	6.9%	62.3%	30.8%	18.3%
Matryoshka	256	3.32	48	d_2 -favouring	-0.796	26.6%	50.6%	22.8%	53.8%
			62	d_1 -favouring	+0.591	6.0%	71.2%	22.8%	20.8%
	192	3.98	30	d_1 -favouring	+0.866	72.5%	16.5%	11.0%	86.8%

E. Further Generalisation Experiment Details

To test whether special latents can be steered to swap outputs in other SAEs, we provide the following pipeline that works given any SAE, which is available in our public codebase:

1. Store the SAE activations for all input bigrams by running a single forward pass over the full dataset and recording each SAE latent’s activation magnitude.
2. Identify special latents using $\Delta\alpha$ correlation $|r| > 0.3$ as a heuristic, as described previously.
3. For each identified special latent, run a batched steering grid search over all input bigrams. For each bigram and each scale factor $p \in [0, 10]$ with step size 0.05 (201 values), record the logits at output positions o_1 and o_2 under the steered reconstruction $\hat{\mathbf{r}}_s^{L_1}$. Bigrams for which the latent does not activate ($m = 0$) are skipped immediately since scaling has no effect.
4. Find crossover scales for each bigram:
 - For o_1 : since o_1 logits are empirically linear in p , fit lines to $\ell_{d_1}(p)$ and $\ell_{d_2}(p)$ and compute the analytical intersection. Inputs for which either logit series fails an $R^2 \geq 0.999$ linearity check, or for which the intersection falls on a negative scale, are flagged and excluded.




Figure 12. These strip plots show how the number of special latents (identified by the latent activation’s correlation with $\Delta\alpha$ using threshold $|r| > 0.3$, as in Section 3.1) differs across all 2,613 trained SAEs from the sweep in Table 4. **Left:** A large majority of the SAEs achieve $L0 \approx 3$. However, only those with special latents achieve high performance. **Right:** SAEs with no special latents tend to perform much worse than those with them. Most SAEs learn two special latents.




Figure 13. These strip plots show how the number of special latents (identified by the latent activation’s correlation with $\Delta\alpha$ using threshold $|r| > 0.3$, as in Section 3.1) differs across all **BatchTopK** SAEs from the sweep in Table 4.



Figure 14. These strip plots show how the number of special latents (identified by the latent activation’s correlation with $\Delta\alpha$ using threshold $|r| > 0.3$, as in Section 3.1) differs across all **JumpReLU** SAEs from the sweep in Table 4.




Figure 15. These strip plots show how the number of special latents (identified by the latent activation’s correlation with $\Delta\alpha$ using threshold $|r| > 0.3$, as in Section 3.1) differs across all **Matryoshka** SAEs from the sweep in Table 4.

- For o_2 : detect sign changes in $\ell_{d_1}(p) - \ell_{d_2}(p)$ on the coarse $[0, 10]$ grid, then refine each crossing to three decimal places using bisection search. All bisection tasks across the batch are executed in parallel to avoid sequential per-input forward passes.

5. Determine a valid swap zone $[p_{lo}, p_{hi}]$ per bigram by intersecting the o_1 and o_2 crossover constraints with the scale ranges over which $\text{argmax}(o_1) = d_2$ and $\text{argmax}(o_2) = d_1$ hold simultaneously. If no contiguous window satisfying both conditions exists, the bigram is recorded as a failure with a specific reason (see Section 3.2).
6. Verify each swap zone by running the model at the midpoint $p^* = (p_{lo} + p_{hi})/2$ and confirming that the model output is (d_2, d_1) .
7. Report the fraction of bigrams for which a valid swap zone was found and the swap verified, along with a breakdown of failure reasons for the remainder.

Steps 3-5 of the pipeline are very computationally expensive, so we run the steering analysis on six selected high-performing SAEs (Table 8) rather than on all SAEs in Figure 12. These checkpoints were chosen to include two BatchTopK, two JumpReLU and two Matryoshka SAEs while varying dictionary size and sparsity. The results are shown in Table 10.

In future work, this study should be extended to further configurations for robust statistical significance.

F. Limitations

- A. **Minimal toy model.** The transformer is trained on a single narrow task and omits components standard in pretrained LLMs (MLPs, LayerNorm, multi-head attention, value and output projections), with a vocabulary of only 100 symbols. Whether relative magnitude-based composition appears in less constrained models remains an open question.
- B. **Single-latent steering.** The steering pipeline handles one special latent at a time. The co-activation experiments in Appendix D suggest that simultaneous multi-latent steering could recover additional swaps.
- C. **SAE selection bias.** Our primary SAE was selected for having few dead latents, biasing toward high-quality SAEs. A fully representative study would require analysis across a random sample of all trained SAEs, not just well-performing ones.
- D. **Limited SAE steering generalisation.** Special latents emerge across BatchTopK, JumpReLU, and Matryoshka architectures in the broader correlation analysis, but steering was tested on only six selected high-performing checkpoints. Broader statistical significance testing across more configurations is needed.